

Algorithm 892: DISPMODULE, a Fortran 95 Module for Pretty-Printing Matrices

KRISTJAN JONASSON
University of Iceland

A standard Fortran 95 module for printing scalars, vectors, and matrices to external files is provided. The module can display variables of default kind of all intrinsic types (integer, real, complex, logical, and character), and add-on modules are provided for data of the nondefault kind. The main module is self-contained and incorporating it only requires that it be compiled and linked with a program containing a “use dispmodule” statement. A generic interface and optional parameters are used, so that the same subroutine name, DISP, is used to display items of different data type and rank, irrespective of display options. The subroutine is quite versatile, and hopefully can improve Fortran’s competitiveness against other array programming languages. The module also contains a function TOSTRING to convert numerical scalars and vectors to strings.

Categories and Subject Descriptors: D.2.2 [Software Engineering]: Design Tools and Techniques—*Software libraries*; D.2.5 [Software Engineering]: Testing and Debugging—*Debugging aids*; D.4.4 [Operating Systems]: Communications Management—*Input/output*; D.3 [Programming Languages]; E.1 [Data Structures]—*Arrays*; G.4 [Mathematical Software]—*User interfaces*; H.5.2 [Information Interfaces and Presentation]: User Interfaces—*Prototyping, screen design (e.g., text, graphics, color), user-centered design*

General Terms: Algorithms, Design

Additional Key Words and Phrases: Fortran 95, matrix pretty-printing, matrix printing, output utilities, array programming language

ACM Reference Format:

Jonasson, K. 2009. Algorithm 892: DISPMODULE, a Fortran 95 module for pretty-printing matrices. ACM Trans. Math. Softw. 36, 1, Article 6 (March 2009), 7 pages. DOI = 10.1145/1486525.1486531 <http://doi.acm.org/10.1145/1486525.1486531>

1. INTRODUCTION

One of the reasons for the great popularity of numerical computation systems such as Matlab [Moler 2004], S-plus [Krause and Olson 2000], and their open source clones (e.g., Octave [Eaton 2002], Scilab [Mrkaic 2001], and R [Chambers

Author’s address: Department of Computer Science, School of Science and Engineering, University of Iceland, Hjardarhaga 4, 107 Reykjavik, Iceland; email: jonasson@hi.is.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org.
© 2009 ACM 0098-3500/2009/03-ART6 \$5.00 DOI 10.1145/1486525.1486531 <http://doi.acm.org/10.1145/1486525.1486531>

ACM Transactions on Mathematical Software, Vol. 36, No. 1, Article 6, Publication date: March 2009.

2007]) stems from the ability of these systems to treat vectors and matrices as fundamental data types. At the heart of these systems are languages that may be classified as *array programming languages*. The first array programming language, APL, was developed in 1957 and it, along with its successor, J, still enjoy considerable popularity [Iverson 1991]. A feature of all these languages is that the syntax for displaying arrays in traditional mathematical format is very simple. As an example, the Matlab command `disp(A)` will pretty-print the matrix `A` using a fairly versatile default format. The Matlab commands `A = exp([3,2;-3,-2]); disp(A)` will display:

```
20.0855    7.3891
 0.0498    0.1353
```

It may be argued that the array programming features account for the widespread use of these languages for writing prototype programs, which are later replaced by programs written in compiled high-level languages such as C or Fortran.

In the 1990s, array programming in Fortran became possible with compilers for Fortran 90 and Fortran 95 [ISO/IEC 1997]. The future of Fortran is in many respects bright [Reid 2003, 2006; Metcalf et al. 2003; ISO/IEC 2004]. However, even with recent enhancements Fortran still lacks the ability to print matrices with a concise command. To display a general matrix `A`, one needs a sequence of statements such as

```
do i=1,size(A,1)
  write(*,'(20G12.4)') A(i,:)
end do
```

which for the matrix given above will produce a rather misaligned output:

```
20.09          7.389
0.4979E-01     0.1353
```

Using an F edit descriptor, an aligned display similar to the one given by Matlab is obtained, but this comes at a cost, as matrices with large elements can no longer be displayed. Fortran also offers list format for writing with default format. To take an example, the statement “`write(*,*) A`” might display 20.085537 0.049787067 7.389056 0.13533528. However, aligned output of matrices with multiple rows is not possible using this, and in addition the result is compiler dependent.

The purpose of the current algorithm is to remedy this deficiency, by providing a module written in standard Fortran, that will, with minimal fuss, display a vector or a matrix with sensible default format. At the same time, the module offers considerable flexibility in controlling the format of the output (much more than Matlab’s `disp`). It is possible to specify the edit descriptor, to label the displayed items, to number the rows and columns, and to direct the output to a file, to the screen, or even (when calling Fortran from Matlab) to the Matlab command window. Several matrices/vectors may be written side by side.

The algorithm is especially useful for debugging purposes, and for the preliminary display of numerical results, but it is general enough that it may in

many cases be used for the final display of such results. It uses a generic interface and has optional parameters, so that the same subroutine name, `DISP`, is used to display items of different data types and ranks, with or without labels, and using default or specified format (in Fortran terminology, scalars have rank 0, vectors rank 1, matrices rank 2, etc.). All the intrinsic data types of Fortran are supported, that is, integer, real (both single and double precision), complex, logical, and character. The module is accompanied by documentation in a comprehensive user manual, in four different formats (plain text, html, pdf, and Microsoft Word).

Among related earlier work, one can mention the Fortran 95 package `Matran` [Stewart 2003], which contains a pretty-print subroutine for items of type `Matran` matrix, but that routine is of much more limited scope than the present one. More versatile matrix printing routines are in the `NAG Libraries` [NAG 2000, 2006], but these are not free, and even if they have a few features that are absent in the current algorithm, the latter is in most respects more flexible.

2. ALGORITHM COMPONENTS

The present algorithm consists of one principal Fortran 95 module, `DISPMODULE`, which handles the display of data types of default kind (that are guaranteed to exist by the Fortran standard), and several add-on modules to handle data types of nondefault kind (such as byte integers and quadruple precision reals). The main procedures are the subroutine `DISP`, which is used to display a scalar, vector, or matrix, and the function `TOSTRING`, which is used to change scalars or vectors into character strings. In addition, there are four subroutines to control settings for how the data is edited and one function to retrieve settings for `DISP`.

2.1 Subroutine `DISP`

Simple ways to call the subroutine `DISP` are

```
CALL DISP(X)
CALL DISP(TITLE, X)
CALL DISP(X, FMT)
CALL DISP(TITLE, X, FMT)
```

where `x` is the item to be displayed, `TITLE` is a character string used to label the displayed item, and `FMT` is an edit descriptor to control the display. Assume that $a_{ij} = \exp(i + j - 1)$ and $b_{ij} = \exp(i^j)$, $i, j = 1, \dots, 3$. Then examples of calls are `CALL DISP('A = ', A)`, which will write out

```
A =  2.718   7.389  20.086
      7.389  20.086  54.598
      20.086  54.598 148.410
```

`CALL DISP(B)`, which displays

```
2.71828E+00  2.71828E+00  2.71828E+00
7.38906E+00  5.45981E+01  2.98096E+03
2.00855E+01  8.10308E+03  5.32048E+11
```

and `CALL DISP(A(1:2,:), 'F0.5')`, which displays

```
2.71828  7.38906  20.08554
7.38906  20.08554  54.59815.
```

A call with a complete list of arguments is `CALL DISP (TITLE, X, FMT, FMT_IMAG, ADVANCE, DIGMAX, LBOUND, ORIENT, SEP, STYLE, TRIM, UNIT, ZEROAS)`. All the arguments are optional, and the purpose, data type, and possible values of each argument are given in Table 1 in the user manual. Not all arguments are compatible with all data types or ranks of `x`. It is, for instance, only possible to specify `FMT_IMAG` for complex `x` and `ORIENT` for vectors. Argument association for arguments after (or starting with) `FMT` will usually be realized with argument keywords, for example, `CALL DISP('A=', A, UNIT=3, ZEROAS='.'')`. Examples demonstrating some of the possibilities are:

```
CALL DISP('MATRIX', A, STYLE='UNDERLINE & NUMBER', UNIT=8, DIGMAX=4)
```

which with the matrix `A` given above will send the following to a formatted file on unit 8,

```

                MATRIX
-----
              1      2      3
1      2.7    7.4    20.1
2      7.4    20.1    54.6
3     20.1    54.6   148.4
```

and

```
K = [-3,0,12,14,0]
CALL DISP('K', K, STYLE='PAD', ORIENT='ROW', SEP=' ', ZEROAS='.'')
```

which will display on the asterisk unit (usually the screen) as

```

-----K-----
-3 . 12 14 .
```

Further examples and details can be found in the user manual.

2.2 The Function TOSTRING

Many programming languages have built-in functions that change numbers to strings. It is, for instance, possible with Java to write

```
System.out.println("The square of " + x + " is " + x*x);
```

(Java also has a function, `Float.toString` that may be used instead). In Matlab, one may write

```
disp(['The square of ' num2str (x) ' is ' num2str (x*x)])
```

If `x` is 1.5, both commands display “The square of 1.5 is 2.25.” The function `TOSTRING` offers similar functionality. The statement

```
call disp('The square of ' // tostring(x) // ' is ' // tostring(x*x))
```

will produce the same output as the Java and Matlab commands. As shown in the user manual, it is possible to achieve a similar effect in Fortran using internal files and list-directed output, but this is cumbersome and the result is compiler dependent. TOSTRING can change numeric and logical scalars and vectors into character strings, and it is possible to specify how to format the string. An example of using TOSTRING is the following creation of variable format:

```
fmt = '(F' // tostring(w) // '.' // tostring(d) // ')'
write(*,fmt) A
```

For further examples and description, see the user manual.

2.3 Changing and Retrieving Settings for DISP and TOSTRING

The default settings used by DISP may be changed with the subroutine DISP_SET, which with a complete list of arguments may be called with

```
CALL DISP_SET (ADVANCE, DIGMAX, MATSEP, ORIENT, SEP, STYLE, UNIT, ZEROAS)
```

All the arguments are optional, and the effect of calling DISP_SET with a specified argument present is to change the default value for the corresponding argument of DISP. For instance, after calling CALL DISP_SET(UNIT=7), subsequent DISP-calls will send output to a file connected to unit 7. The only argument that is unique to DISP_SET is MATSEP (for *matrix separator*), which is a character string that is used to separate items displayed with ADVANCE='NO' (the default separator is a string with three spaces). To take an example, the sequence of calls

```
CALL DISP_SET (MATSEP = ' | ')
CALL DISP([11,12,13], ADVANCE = 'NO')
CALL DISP([.TRUE., .FALSE., .TRUE.], ADVANCE = 'NO')
CALL DISP(['A', 'B', 'C'])
```

will display

```
11 | T | A
12 | F | B
13 | T | C
```

Similarly, there is a function TOSTRING_SET to change default settings for TOSTRING. There are also functions to return the settings to the original (or *factory*) defaults. These are called DISP_SET_FACTORY and TOSTRING_SET_FACTORY. The current default settings for DISP may be retrieved with the function DISP_GET, which returns a structure (say DS) of type DISP_SETTINGS (declared in DISPMODULE) with these settings. The settings may be reapplied to DISP by calling DISP_SET(DS). The user manual contains details.

2.4 Testing

The algorithm contains a test program, TEST_DISPMODULE, which applies modern testing techniques to ascertain the correctness of DISPMODULE. The test program

calls all the procedures of the module, and all arguments are tried. Testing of all possible data types for x is allowed for. For arguments that have a limited range of allowed values, all these values are tried. In addition, all the examples given in the user manual and in this article are tested. It is, however, impossible to test all possible combinations of x data type, arguments, and argument values, so one must be content with a partial test, accompanied by the examination of the source code of the module(s) and the test program.

There is a parameter declared at the beginning of the test program to control the verbosity level, or the amount of output. It has three possible settings: minimal output, a short report, and a detailed report (including the result of all `DISP` and `TOSTRING` invocations). In all cases the report ends with “OK” if all tests are passed. Further instructions for operation of the test program are given in the user manual.

The algorithm has been tested successfully with recent versions of the following compilers: under Microsoft Windows: gfortran, g95, f95 from NAG, and ifort from Intel; under Ubuntu Linux: gfortran and g95; under Sun Solaris: g95; under Suse Linux: ifort; under Fedora Linux: ifort, af95 from Absoft, and pathf95 from Pathscale. It passes tests of g95, gfortran, and ifort of conforming to the Fortran 95 standard.

2.5 Algorithm Limitations and Assumptions on the Compiler

Among major limitations of the module is that `DISP` cannot display very wide matrices, as there is no provision for breaking matrices (or matrix rows) vertically. The compilers tried can, however, all handle record lengths of 1024 characters or more, and since the output of `DISP` is intended for human reading this should not be too serious a drawback. Another important limitation is that displaying arrays of rank ≥ 3 is not supported. Some minor limitations are that `zeroas` is not supported for complex variables, `TOSTRING` cannot handle character variables, and 2-byte logicals are not directly supported.

Among the things which the NAG matrix printing routines can do, and `DISP` cannot, is writing matrices with special storage schemes, such as triangular, band, or symmetric. Another ability of the NAG routines missing from `DISP` is specifying character strings for matrix row and column labels.

An assumption on the compiler in the design of the algorithm is that, if A and B are real variables with $|A| < |B|$, and writing B with *Fw.d* editing does not produce w asterisks, then neither does writing A . One possible problem that hopefully does not occur involves compilers that automatically prepend a plus sign when writing positive numbers. Every effort has been made to use `SS` editing to suppress potential plus signs, but no compiler that defaults to `SP` editing is known to the author.

The feature of writing matrices side by side (with `ADVANCE = 'NO'`) is implemented by allocating memory for the blocks waiting to be written out. To ensure that this memory is freed, the last call of `DISP` on each unit should have `ADVANCE = 'YES'` in effect.

REFERENCES

- CHAMBERS, J. M. 2007. *Software for Data Analysis: Programming with R*. Springer-Verlag, New York, NY.
- EATON, J. W. 2002. *GNU Octave Manual*. Network Theory Limited, Bristol, U.K.
- ISO/IEC. 1997. Information technology—programming languages—Fortran-part 1: Base language. ISO/IEC Tech. rep. 1539-1:1997. ISO, Geneva, Switzerland.
- ISO/IEC. 2004. Information technology—programming languages—Fortran-part 1: Base language. ISO/IEC Tech. rep. 1539-1-2004. ISO, Geneva.
- IVERSON, K. E. 1991. A personal view of APL. *IBM Syst. J.* 30, 4, 582–593.
- KRAUSE, A. AND OLSON, M. 2000. *The Basics of S and S-PLUS (second ed.)*, Springer-Verlag, New York.
- METCALF, M., REID, J. K. AND COHEN, M. 2004. *Fortran 95/2003 Explained*. Oxford University Press, Oxford, U.K.
- MOLER, C. 2004. *Numerical Computing with MATLAB*, SIAM, Philadelphia, PA.
- MRKAIC, M. 2001. Scilab as an econometric programming system. *J. Appl. Econometr. Systems* 16, 4, 553–559.
- NAG. 2000. *NAG Fortran 90 Library Manual*. Numerical Algorithms Group, Oxford, U.K.
- NAG. 2006. *NAG Fortran Library Manual*. Numerical Algorithms Group, Oxford, U.K.
- REID, J. K. 2003. The future of Fortran. *Comput. Sci. Eng.* 5, 4, 59–67.
- REID, J. K. 2006. Fortran is getting more and more powerful. In *Applied Parallel Computing, State of the Art in Scientific Computing*, J. Dongarra, K. Madsen, and J. Wasniewski, Eds. Lecture Notes in Computer Science, vol. 3732. Springer, Berlin, Germany, 33–42.
- STEWART, G. W. 2003. *MATRAN: A Fortran-95 Matrix Wrapper*. Tech. rep. 4522. University of Maryland, Department of Computer Science, College Park, MD.

Received April 2008; accepted August 2008